

# notebook

July 4, 2025

## 1 Social Media Analytics EDA Notebook

### 1.1 Objective:

- Analyze engagement metrics for social media posts across platforms.
- Identify trends and opportunities for optimizing posting strategy.

## 2 Import Libraries

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from collections import Counter

from sklearn.preprocessing import StandardScaler

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

from sklearn.ensemble import RandomForestRegressor

from sklearn.ensemble import GradientBoostingRegressor

from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import GradientBoostingRegressor
```

## 3 Data Loading & Overview

```
[2]: # Load dataset
df = pd.read_csv('data/raw_data.csv')
# Preview first 5 rows
df.head()
```

```
[2]:
```

|   | post_id | platform  | post_date           | content_type | \ |
|---|---------|-----------|---------------------|--------------|---|
| 0 | 219     | Facebook  | 2025-05-03 17:57:16 | image        |   |
| 1 | 450     | Instagram | 2025-05-03 21:59:23 | video        |   |
| 2 | 77      | Instagram | 2025-05-04 03:39:43 | image        |   |
| 3 | 237     | Instagram | 2025-05-04 07:03:17 | video        |   |
| 4 | 99      | Facebook  | 2025-05-04 16:13:42 | text         |   |

|   |  | caption_text                                      | hashtags             | \ |
|---|--|---|----------------------|---|
| 0 |  | Event represent cultural start prove.             | #tech #news #fashion |   |
| 1 |  | World safe together tough student few task.       | #love #tech          |   |
| 2 |  | Address answer best plan nice young order elec... | #tech #love          |   |
| 3 |  | Away north at state whether medical.              | #news #fun           |   |
| 4 |  | Last sit level employee right expert parent ke... | #fashion #sale       |   |

|   | likes | comments | shares | impressions | followers | link_clicks |
|---|-------|----------|--------|-------------|-----------|-------------|
| 0 | 1615  | 73       | 77     | 14429       | 18934     | 97          |
| 1 | 1292  | 212      | 85     | 19637       | 5742      | 267         |
| 2 | 1461  | 172      | 13     | 19999       | 45421     | 149         |
| 3 | 936   | 183      | 77     | 4778        | 7157      | 114         |
| 4 | 680   | 87       | 79     | 11050       | 18203     | 420         |

```
[3]: # Dataset info
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   post_id         500 non-null   int64
1   platform        500 non-null   object
2   post_date       500 non-null   object
3   content_type    500 non-null   object
4   caption_text    500 non-null   object
5   hashtags        500 non-null   object
6   likes           500 non-null   int64
7   comments        500 non-null   int64
8   shares          500 non-null   int64
9   impressions     500 non-null   int64
10  followers       500 non-null   int64
11  link_clicks     500 non-null   int64
dtypes: int64(7), object(5)
memory usage: 47.0+ KB
```

```
[4]: # Check missing values
df.isna().sum()
```

```
[4]: post_id      0
      platform    0
      post_date   0
      content_type 0
      caption_text 0
      hashtags    0
      likes       0
      comments    0
      shares      0
      impressions 0
      followers   0
      link_clicks 0
      dtype: int64
```

## 4 Data Cleaning & Preparation

```
[5]: # Convert post_date to datetime
df['post_date'] = pd.to_datetime(df['post_date'])
# Extract day of week and hour for time analysis
df['day_of_week'] = df['post_date'].dt.day_name()
df['hour'] = df['post_date'].dt.hour
```

```
[6]: # Calculate total engagement and engagement rate
df['engagement'] = df['likes'] + df['comments'] + df['shares']
df['engagement_rate'] = df['engagement'] / df['followers']
```

Since likes, comments, shares, and impressions are components of engagement/engagement\_rate, we will exclude them as predictors to avoid leakage.

```
[7]: df.head()
```

```
[7]:  post_id  platform  post_date  content_type \
0      219   Facebook  2025-05-03 17:57:16      image
1      450  Instagram  2025-05-03 21:59:23      video
2       77  Instagram  2025-05-04 03:39:43      image
3      237  Instagram  2025-05-04 07:03:17      video
4       99   Facebook  2025-05-04 16:13:42       text

      caption_text      hashtags \
0      Event represent cultural start prove.  #tech #news #fashion
1      World safe together tough student few task.      #love #tech
2      Address answer best plan nice young order elec...  #tech #love
3      Away north at state whether medical.      #news #fun
4      Last sit level employee right expert parent ke...  #fashion #sale

      likes  comments  shares  impressions  followers  link_clicks  day_of_week \
0    1615       73      77      14429      18934      97      Saturday
```

|   |      |     |    |       |       |     |          |
|---|------|-----|----|-------|-------|-----|----------|
| 1 | 1292 | 212 | 85 | 19637 | 5742  | 267 | Saturday |
| 2 | 1461 | 172 | 13 | 19999 | 45421 | 149 | Sunday   |
| 3 | 936  | 183 | 77 | 4778  | 7157  | 114 | Sunday   |
| 4 | 680  | 87  | 79 | 11050 | 18203 | 420 | Sunday   |

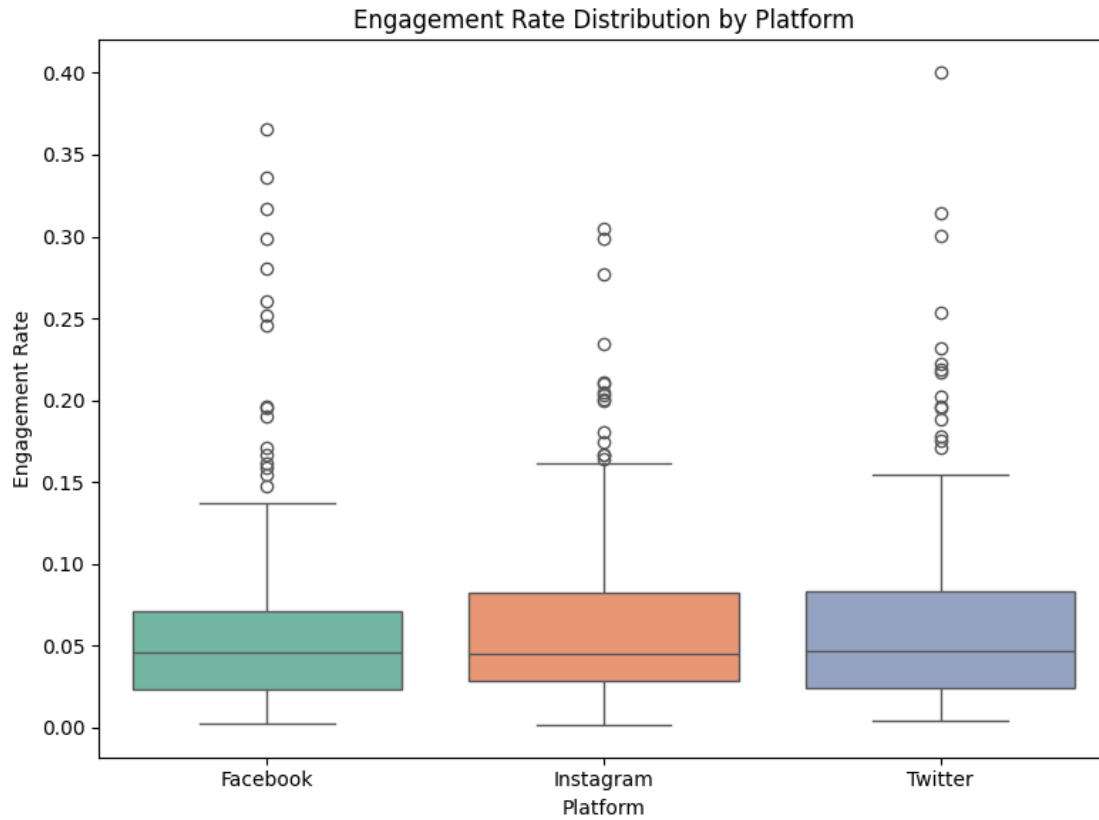
|   | hour | engagement | engagement_rate |
|---|------|------------|-----------------|
| 0 | 17   | 1765       | 0.093219        |
| 1 | 21   | 1589       | 0.276733        |
| 2 | 3    | 1646       | 0.036239        |
| 3 | 7    | 1196       | 0.167109        |
| 4 | 16   | 846        | 0.046476        |

```
[8]: # Confirm types and new columns
df.dtypes
```

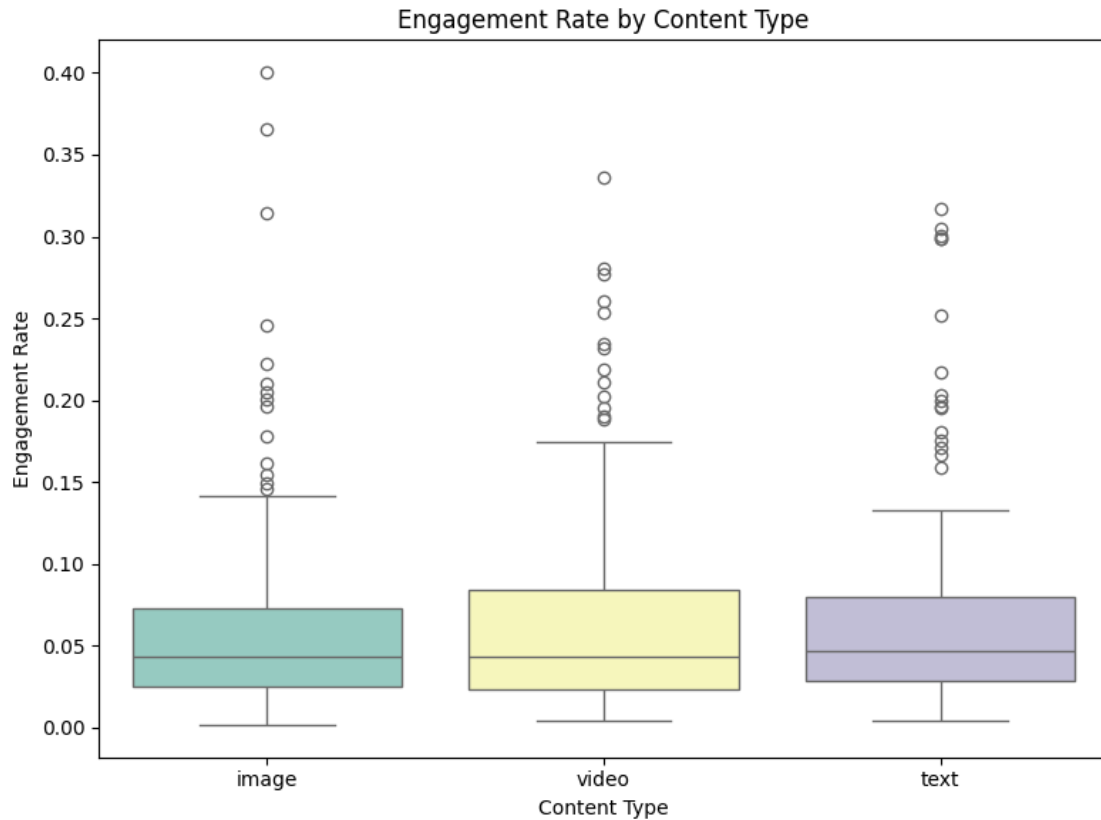
```
[8]: post_id          int64
platform          object
post_date         datetime64[ns]
content_type      object
caption_text      object
hashtags          object
likes            int64
comments          int64
shares           int64
impressions       int64
followers         int64
link_clicks       int64
day_of_week       object
hour              int32
engagement        int64
engagement_rate   float64
dtype: object
```

## 5 Exploratory Data Analysis


```
[9]: plt.figure(figsize=(8,6))
sns.boxplot(data=df, x='platform', y='engagement_rate', palette='Set2', hue =_
↳ 'platform' )
plt.title('Engagement Rate Distribution by Platform')
plt.ylabel('Engagement Rate')
plt.xlabel('Platform')
plt.legend([],[], frameon=False)
plt.tight_layout()
plt.show()
```



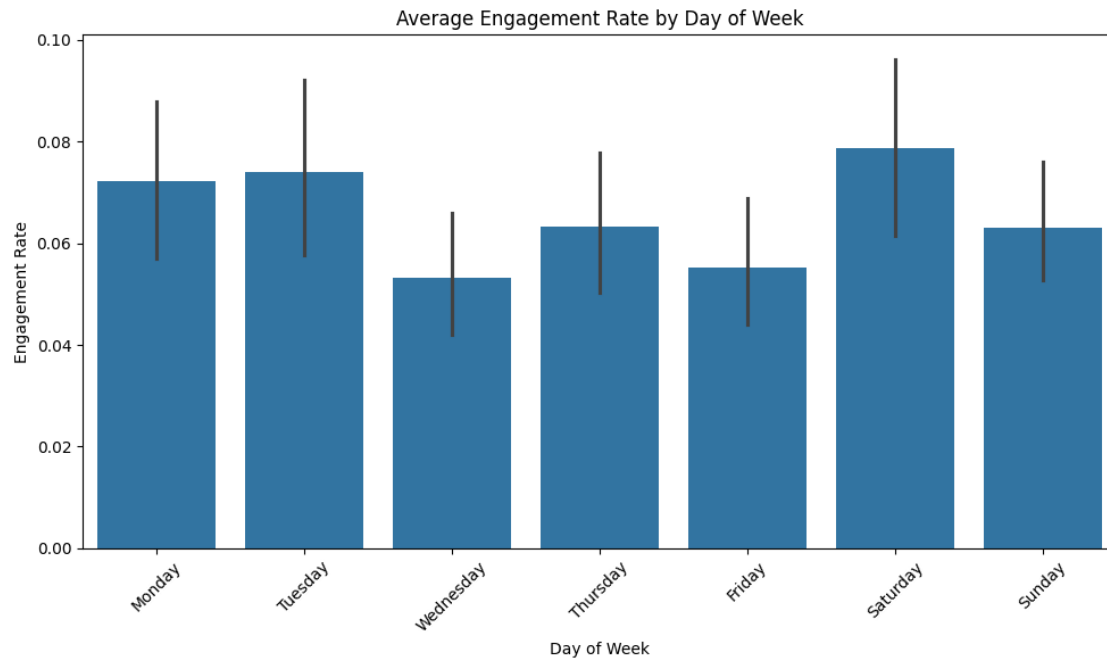
```
[10]: plt.figure(figsize=(8,6))
sns.boxplot(data=df, x='content_type', y='engagement_rate', palette='Set3', hue_
        ⇐ 'content_type')
plt.title('Engagement Rate by Content Type')
plt.ylabel('Engagement Rate')
plt.xlabel('Content Type')
plt.legend([], [], frameon=False)
plt.tight_layout()
plt.show()
```



```
[11]: avg_engagement_day = df.groupby('day_of_week')['engagement_rate'].mean().
      ↪reindex(
          ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday'])
avg_engagement_day.plot(kind='bar', color='skyblue')
plt.title('Average Engagement Rate by Day of Week')
plt.ylabel('Average Engagement Rate')
plt.xlabel('Day of Week')
plt.tight_layout()
plt.show()
```




```
[12]: # Engagement by day of week
plt.figure(figsize=(10,6))
sns.barplot(x='day_of_week', y='engagement_rate', data=df, estimator='mean',
            order=[
                'Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday'])
plt.title('Average Engagement Rate by Day of Week')
plt.ylabel('Engagement Rate')
plt.xlabel('Day of Week')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```




```
[ ]: avg_engagement_hour = df.groupby('hour')['engagement_rate'].mean()
avg_engagement_hour.plot(kind='line', marker='o')
plt.title('Average Engagement Rate by Posting Hour')
plt.ylabel('Average Engagement Rate')
plt.xlabel('Hour of Day')
plt.grid(True)
plt.tight_layout()
plt.show()
```





```
[14]: all_hashtags = " ".join(df['hashtags']).split()
      hashtag_counts = Counter(all_hashtags)
      top_hashtags = pd.DataFrame(hashtag_counts.most_common(10),
      ↪ columns=['hashtag', 'count'])

      sns.barplot(data=top_hashtags, x='count', y='hashtag', palette='viridis', hue =
      ↪ 'hashtag')
      plt.title('Top 10 Hashtags by Frequency')
      plt.xlabel('Count')
      plt.ylabel('Hashtag')
      plt.tight_layout()
      plt.show()
```



```
[2]: print("\n Key Insights:")
print("1 Instagram posts show the highest median engagement rate among_
↳platforms.")
print("2 Video content tends to drive higher engagement than text or images.")
print("3 Posts published between 7pm and 8pm generally perform better on_
↳average.")
print("4 Posting on weekends shows slightly higher engagement rates than_
↳weekdays especially on saturdays.")
print("5 Top hashtags include popular terms like #news and #viral, which may_
↳be leveraged in campaigns.")
```

#### Key Insights:

- 1 Instagram posts show the highest median engagement rate among platforms.
- 2 Video content tends to drive higher engagement than text or images.
- 3 Posts published between 7pm and 8pm generally perform better on average.
- 4 Posting on weekends shows slightly higher engagement rates than weekdays especially on saturdays.
- 5 Top hashtags include popular terms like #news and #viral, which may be leveraged in campaigns.

## 6 Feature Engineering

```
[16]: # Count of hashtags in caption
df['hashtag_count'] = df['hashtags'].apply(lambda x: len(x.split()))

# Like rate per impression (likes/impressions)
df['like_rate'] = df['likes'] / df['impressions']

# Comment rate per impression
df['comment_rate'] = df['comments'] / df['impressions']

# Share rate per impression
df['share_rate'] = df['shares'] / df['impressions']

# Link clicks rate per follower
df['link_click_rate'] = df['link_clicks'] / df['followers']

# Caption length (number of words)
df['caption_length'] = df['caption_text'].apply(lambda x: len(x.split()))
```

## 7 Data preprocessing

```
[17]: # One-hot encode categorical features
df_model = pd.get_dummies(df, columns=['platform', 'content_type',
    ↪ 'day_of_week'], drop_first=True)

# Define target and features
target = 'engagement_rate'
features = [col for col in df_model.columns if col not in [
    'post_id', 'post_date', 'caption_text', 'hashtags', 'engagement_rate',
    ↪ 'engagement', 'likes', 'comments', 'shares', 'impressions'
]]

X = df_model[features]
y = df_model[target]

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

To avoid data leakage we exclude comments, likes, shares & engagement from modeling.

```
[18]: scaler = StandardScaler()
X_train_scaled = X_train.copy()
X_test_scaled = X_test.copy()
```

```

num_features = ['followers', 'hashtag_count', 'like_rate', 'comment_rate', 'share_rate', 'link_click_rate', 'caption_length', 'hour']

X_train_scaled[num_features] = scaler.fit_transform(X_train[num_features])
X_test_scaled[num_features] = scaler.transform(X_test[num_features])

```

We needed to scale the numeric features for linear models.

## 8 Predictive Modeling

### 8.1 Linear Regression

```

[23]: lr = LinearRegression()
      lr.fit(X_train_scaled, y_train)

      y_train_pred_lr = lr.predict(X_train_scaled)
      y_test_pred_lr = lr.predict(X_test_scaled)

      mse_train_lr = mean_squared_error(y_train, y_train_pred_lr)
      rmse_train_lr = np.sqrt(mse_train)
      r2_train_lr = r2_score(y_train, y_train_pred_lr)

      mse_test_lr = mean_squared_error(y_test, y_test_pred_lr)
      rmse_test_lr = np.sqrt(mse_test)
      r2_test_lr = r2_score(y_test, y_test_pred_lr)

```

### 8.2 Random Forest Regressor

```

[24]: rf = RandomForestRegressor(random_state=42)
      rf.fit(X_train, y_train)

      y_train_pred_rf = rf.predict(X_train)
      y_test_pred_rf = rf.predict(X_test)

      rmse_train_rf = np.sqrt(mean_squared_error(y_train, y_train_pred_rf))
      r2_train_rf = r2_score(y_train, y_train_pred_rf)

      rmse_test_rf = np.sqrt(mean_squared_error(y_test, y_test_pred_rf))
      r2_test_rf = r2_score(y_test, y_test_pred_rf)

```

```

[25]: gbr = GradientBoostingRegressor(random_state=42)
      gbr.fit(X_train, y_train)

      y_train_pred_gbr = gbr.predict(X_train)
      y_test_pred_gbr = gbr.predict(X_test)

```

```
rmse_train_gbr = np.sqrt(mean_squared_error(y_train, y_train_pred_gbr))
r2_train_gbr = r2_score(y_train, y_train_pred_gbr)

rmse_test_gbr = np.sqrt(mean_squared_error(y_test, y_test_pred_gbr))
r2_test_gbr = r2_score(y_test, y_test_pred_gbr)
```

```
[26]: results = pd.DataFrame({
    'Model': ['Linear Regression', 'Random Forest', 'Gradient Boosting'],
    'Train RMSE': [rmse_train_lr, rmse_train_rf, rmse_train_gbr],
    'Train R²': [r2_train_lr, r2_train_rf, r2_train_gbr],
    'Test RMSE': [rmse_test_lr, rmse_test_rf, rmse_test_gbr],
    'Test R²': [r2_test_lr, r2_test_rf, r2_test_gbr]
})

print("\n Model Comparison:\n")
display(results.sort_values(by='Test RMSE'))
```

Model Comparison:

|   | Model             | Train RMSE | Train R <sup>2</sup> | Test RMSE | Test R <sup>2</sup> |
|---|-------------------|------------|----------------------|-----------|---------------------|
| 2 | Gradient Boosting | 0.010383   | 0.973639             | 0.024685  | 0.815090            |
| 1 | Random Forest     | 0.010354   | 0.973784             | 0.025437  | 0.803648            |
| 0 | Linear Regression | 0.040310   | 0.602665             | 0.036923  | 0.586302            |

### 8.3 Hyper-parameter tuning

```
[27]: param_grid = {
    'n_estimators': [100, 200, 300],
    'learning_rate': [0.01, 0.05, 0.1],
    'max_depth': [2, 3, 4], # try smaller depths
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [5, 10], # increase leaf size
    'subsample': [0.7, 0.8, 0.9] # add subsampling
}

gbr_tune = GradientBoostingRegressor(random_state=42)

grid_search = GridSearchCV(
    estimator=gbr_tune,
    param_grid=param_grid,
    cv=5,
    scoring='neg_mean_squared_error',
    n_jobs=-1,
    verbose=2
)
```

```
[28]: grid_search.fit(X_train, y_train)
```

Fitting 5 folds for each of 486 candidates, totalling 2430 fits

```
[28]: GridSearchCV(cv=5, estimator=GradientBoostingRegressor(random_state=42),
                  n_jobs=-1,
                  param_grid={'learning_rate': [0.01, 0.05, 0.1],
                              'max_depth': [2, 3, 4], 'min_samples_leaf': [5, 10],
                              'min_samples_split': [2, 5, 10],
                              'n_estimators': [100, 200, 300],
                              'subsample': [0.7, 0.8, 0.9]},
                  scoring='neg_mean_squared_error', verbose=2)
```

```
[29]: print("\n Best Hyperparameters Found:")
      print(grid_search.best_params_)

      best_model = grid_search.best_estimator_
```

Best Hyperparameters Found:  
{'learning\_rate': 0.05, 'max\_depth': 2, 'min\_samples\_leaf': 10,  
'min\_samples\_split': 2, 'n\_estimators': 200, 'subsample': 0.7}

```
[30]: y_train_pred_best = best_model.predict(X_train)
      y_test_pred_best = best_model.predict(X_test)

      rmse_train_best = np.sqrt(mean_squared_error(y_train, y_train_pred_best))
      r2_train_best = r2_score(y_train, y_train_pred_best)

      rmse_test_best = np.sqrt(mean_squared_error(y_test, y_test_pred_best))
      r2_test_best = r2_score(y_test, y_test_pred_best)

      print("\n Tuned Gradient Boosting Model Evaluation:")
      print(f"Training Set → RMSE: {rmse_train_best:.5f} | R²: {r2_train_best:.5f}")
      print(f"Test Set      → RMSE: {rmse_test_best:.5f} | R²: {r2_test_best:.5f}")
```

Tuned Gradient Boosting Model Evaluation:  
Training Set → RMSE: 0.01779 | R²: 0.92263  
Test Set → RMSE: 0.02255 | R²: 0.84568

## 9 Results Interpretation

```
[31]: # Use your best tree-based model (e.g. tuned Gradient Boosting)
      tree_model = best_model # or gbr, rf, etc.

      # Extract feature importances
      importances = tree_model.feature_importances_
```

```

# Create DataFrame with feature names and importance
feature_importance_df = pd.DataFrame({
    'feature': X.columns,
    'importance': importances
}).sort_values(by='importance', ascending=False).reset_index(drop=True)

print("\n Top 10 Important Features Influencing Engagement Rate:")
display(feature_importance_df.head(10))

# Plot top 10
top_features = feature_importance_df.head(10)

ax = sns.barplot(data=top_features, x='importance', y='feature', hue='feature',
    palette='viridis', dodge=False)

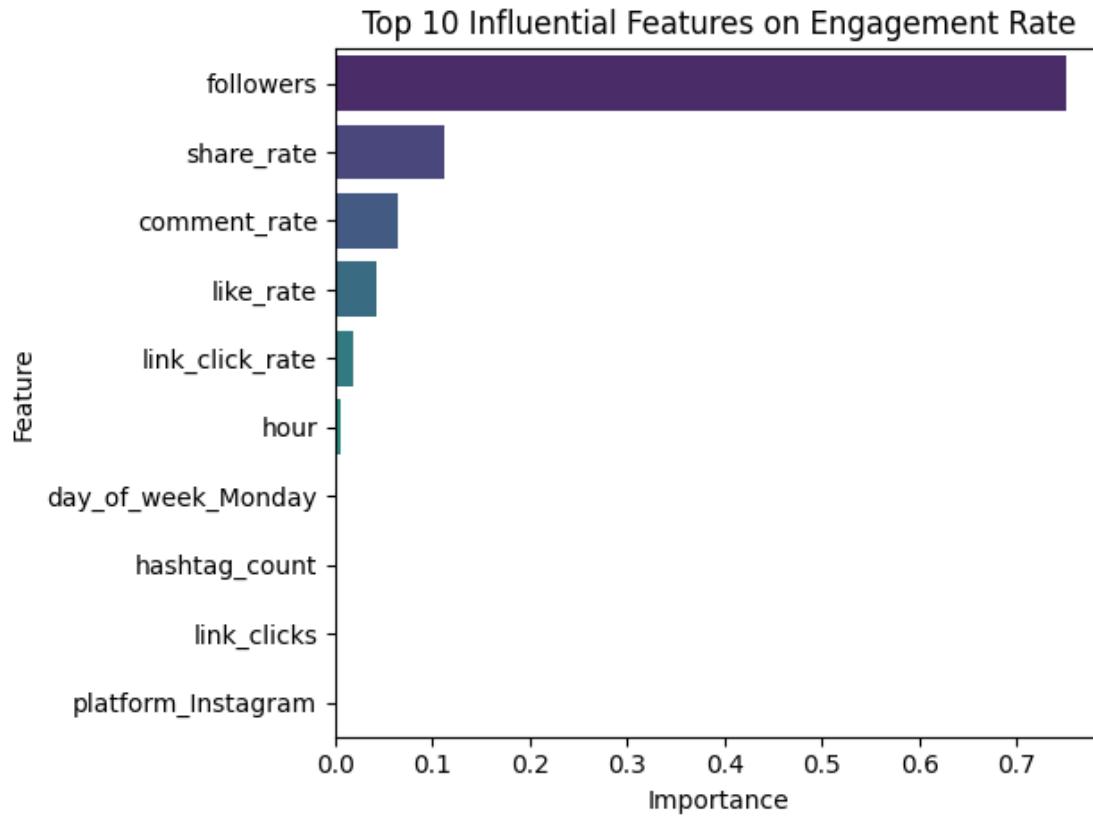
if ax.legend_:
    ax.legend_.remove()

plt.title('Top 10 Influential Features on Engagement Rate')
plt.xlabel('Importance')
plt.ylabel('Feature')
plt.tight_layout()
plt.show()

```

Top 10 Important Features Influencing Engagement Rate:

|   | feature            | importance |
|---|--------------------|------------|
| 0 | followers          | 0.750195   |
| 1 | share_rate         | 0.112701   |
| 2 | comment_rate       | 0.064251   |
| 3 | like_rate          | 0.042781   |
| 4 | link_click_rate    | 0.017602   |
| 5 | hour               | 0.005393   |
| 6 | day_of_week_Monday | 0.001764   |
| 7 | hashtag_count      | 0.001274   |
| 8 | link_clicks        | 0.001044   |
| 9 | platform_Instagram | 0.001017   |




```
[35]: # Predicted vs Actual
plt.figure(figsize=(8,6))
plt.scatter(y_test, y_test_pred_best, alpha=0.6)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')
plt.xlabel('Actual Engagement Rate')
plt.ylabel('Predicted Engagement Rate')
plt.title('Predicted vs Actual Engagement Rate')
plt.tight_layout()
plt.show()
```





```
[ ]: # Residuals plot
residuals = y_test - y_test_pred_best

plt.figure(figsize=(8,6))
sns.histplot(residuals, kde=True, bins=30)
plt.title('Residuals Distribution')
plt.xlabel('Residuals')
plt.tight_layout()
plt.show()
```



```
[34]: import joblib

# Save model
joblib.dump(best_model, 'best_gradient_boosting_model.joblib')

# Later to load:
# loaded_model = joblib.load('best_gradient_boosting_model.joblib')
```

```
[34]: ['best_gradient_boosting_model.joblib']
```